```
import numpy as np
import pandas as pd
import plotly.express as px
import matplotlih.pyplot as plt
import matplotly.graph.objects as go
import warnings
warnings.filterwarnings('ignore')

# Import all data files
users = pd.read_csv('./USER_TAKEHOME.csv')
txns = pd.read_csv('./TRANSACTION_TAKEHOME.csv')
products = pd.read_csv('./PRODUCTS_TAKEHOME.csv')
```

Q1: Are there any data quality issues present?

1. Duplicate Values

Duplicate Products

```
In [2]: # Group same products by BARCODE and count occurrences
        product_counts = products.groupby('BARCODE').size().reset_index(name='COUNT')
        # Filter for duplicate BARCODES (more than one occurrence)
        duplicate_barcodes = product_counts[product_counts['COUNT'] > 1]['BARCODE']
        duplicate_products = products[products['BARCODE'].isin(duplicate_barcodes)]
        duplicate_products = duplicate_products.sort_values(by='BARCODE') #Sorting to visualize same duplicate products together
        print(f"Number of PRODUCT with DUPLICATE BARCODES: {len(duplicate barcodes)}")
        print(f"Number of total unique PRODUCT BARCODES: {products['BARCODE'].nunique()}")
       Number of PRODUCT with DUPLICATE BARCODES: 185
       Number of total unique PRODUCT BARCODES: 841342
In [3]: # Function to highlight both duplicate row grouping and differences
        def highlight_group_differences(df):
            df_highlighted = df.copy()
            df_highlighted['BARCODE'] = df_highlighted['BARCODE'].astype(str)
            # Generate unique colors for each barcode group
            unique_barcodes = df_highlighted['BARCODE'].unique()
            color_palette = px.colors.qualitative.Pastel # Plotly's Pastel color palette
            color_palette_extended = (color_palette * (len(unique_barcodes) // len(color_palette) + 1))[:len(unique_barcodes)]
            barcode_colors = dict(zip(unique_barcodes, color_palette_extended)) # assigning a color to each of the unique barcode
            # Apply background color to entire row based on BARCODE
            row_colors = df_highlighted['BARCODE'].map(barcode_colors)
            row_styles = pd.DataFrame('', index=df.index, columns=df.columns)
            for col in df.columns:
                row_styles[col] = 'background-color: ' + row_colors
            # Compare values within groups to find differences
            for barcode, group in df.groupby('BARCODE'):
                for col in df.columns:
                    if col != 'BARCODE':
                        # Ignore NaN values when checking uniqueness
                        if group[col].dropna().nunique() > 1:
                            row_styles.loc[group.index, col] = 'background-color: yellow'
```

return df_highlighted.style.apply(lambda x: row_styles, axis=None)

Apply the function and display the styled dataframe
highlighted_table = highlight_group_differences(duplicate_products.head(20))
highlighted_table

Out[3]:	CATEGORY_1		CATEGORY_2	CATEGORY_3	CATEGORY_4	MANUFACTURER	BRAND	BARCODE
	349945	Snacks	Candy	Confection Candy	nan	MARS WRIGLEY	STARBURST	400510.0
	99568	Snacks	Candy	Confection Candy	nan	MARS WRIGLEY	STARBURST	400510.0
	139121	Snacks	Candy	Chocolate Candy	nan	PLACEHOLDER MANUFACTURER	BRAND NOT KNOWN	404310.0
	841230	Snacks	Candy	Chocolate Candy	nan	MARS WRIGLEY	M&M'S	404310.0
	274321	Snacks	Crackers	Graham Crackers	nan	TRADER JOE'S	TRADER JOE'S	438711.0
	684662	Snacks	Crackers	Graham Crackers	nan	TRADER JOE'S	TRADER JOE'S	438711.0
	486969	Snacks	Fruit & Vegetable Snacks	Dried Vegetables	nan	TRADER JOE'S	TRADER JOE'S	563178.0
	743615	Snacks	Fruit & Vegetable Snacks	Dried Vegetables	nan	TRADER JOE'S	TRADER JOE'S	563178.0
	269354	Snacks	Nuts & Seeds	Cashews	nan	TRADER JOE'S	TRADER JOE'S	603898.0
	756110	Snacks	Nuts & Seeds	Cashews	nan	TRADER JOE'S	TRADER JOE'S	603898.0
	610681	Snacks	Nuts & Seeds	Snack Seeds	nan	SUNRIDGE FARMS	SUNRIDGE FARMS	701983.0
	645266	Snacks	Chips	Crisps	nan	TRADER JOE'S	TRADER JOE'S	701983.0
	264179	Snacks	Fruit & Vegetable Snacks	Dried Vegetables	nan	TRADER JOE'S	TRADER JOE'S	853743.0
	620981	Snacks	Fruit & Vegetable Snacks	Dried Vegetables	nan	TRADER JOE'S	TRADER JOE'S	853743.0
	368833	Snacks	Dips & Salsa	Hummus	nan	TRADER JOE'S	TRADER JOE'S	906425.0
	325056	Snacks	Dips & Salsa	Hummus	nan	TRADER JOE'S	TRADER JOE'S	906425.0
	87568	Snacks	Chips	Crisps	nan	TRADER JOE'S	TRADER JOE'S	952811.0
	14607	Snacks	Chips	Crisps	nan	TRADER JOE'S	TRADER JOE'S	952811.0
	171015	Snacks	Nuts & Seeds	Covered Nuts	nan	TRADER JOE'S	TRADER JOE'S	969307.0
	681268	Snacks	Nuts & Seeds	Almonds	nan	TRADER JOE'S	TRADER JOE'S	969307.0

ii Insights

- The Products Table contains 185 duplicate products out of a total of 841,342 products
- A product with the same BARCODE appears multiple times with varying values in different fields such as:
 - **BRAND** Different brand names for the same barcode.
 - MANUFACTURER Different manufacturer details for the same barcode.
 - **CATEGORY** Products categorized differently under the same barcode.

Duplicate Transactions

```
In [4]: # Group transactions by RECEIPT_ID and count occurrences
    receipt_counts = txns.groupby('RECEIPT_ID').size().reset_index(name='Count')

# Filter for duplicate RECEIPT_IDs (more than one occurrence)
duplicate_receipts = receipt_counts[receipt_counts['Count'] > 1]

print(f"Number of TXNS with DUPLICATE RECEIPTS: {len(duplicate_receipts)}")
print(f"Number of total unique TXN RECEIPTS: {txns['RECEIPT_ID'].nunique()}")

# Calculate the No of Unique Receipts for each Occurrence Count
occurrence_counts = receipt_counts['Count'].value_counts().reset_index()
occurrence_counts.columns = ['Number of Record in Dataset', 'Number of unique Receipts']
occurrence_counts
```

Number of TXNS with DUPLICATE RECEIPTS: 24440 Number of total unique TXN RECEIPTS: 24440

Out[4]: Number of Record in Dataset Number of unique Receipts

0	2	23920
1	4	488
2	6	26
3	8	5
4	12	1

Receipt ID: 0000d256-4041-4a3e-adc4-5623fb6e0c99

RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FINAL_QUANTITY	FINAL_SALE
0 0000d256-4041-4a3e-adc4-5623fb6e0c99	2024-08-21	2024-08-21 14:19:06.539 Z	WALMART	63b73a7f3d310dceeabd4758	1.530001e+10	1.00	
41567 0000d256-4041-4a3e-adc4-5623fb6e0c99	2024-08-21	2024-08-21 14:19:06.539 Z	WALMART	63b73a7f3d310dceeabd4758	1.530001e+10	1.00	1.54

	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE F	INAL_QUANTITY F	INAL_SALE
1	0001455d-7a92-4a7b-a1d2-c747af1c8fd3	2024-07-20	2024-07-20 09:50:24.206 Z	ALDI	62c08877baa38d1a1f6c211a	NaN	zero	1.49
39291	0001455d-7a92-4a7b-a1d2-c747af1c8fd3	2024-07-20	2024-07-20 09:50:24.206 Z	ALDI	62c08877baa38d1a1f6c211a	NaN	1.00	1.49
Receip [.]	t ID: 00017e0a-7851-42fb-bfab-0baa9	6e23586						
	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCOD	E FINAL_QUANTITY	FINAL_SALE
2	00017e0a-7851-42fb-bfab-0baa96e23586	2024-08-18	2024-08-19 15:38:56.813 Z	WALMART	60842f207ac8b7729e472020	7.874223e+1	0 1.00)
25928	00017e0a-7851-42fb-bfab-0baa96e23586	2024-08-18	2024-08-19 15:38:56.813 Z	WALMART	60842f207ac8b7729e472020	7.874223e+1	0 1.00	2.54
Receip	t ID: 000239aa-3478-453d-801e-66a82	e39c8af						
	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FINAL_QUANTITY	FINAL_SALE
3	000239aa-3478-453d-801e-66a82e39c8af	2024-06-18	2024-06-19 11:03:37.468 Z	FOOD LION	63fcd7cea4f8442c3386b589	7.833997e+11	zero	3.49
41475	000239aa-3478-453d-801e-66a82e39c8af	2024-06-18	2024-06-19 11:03:37.468 Z	FOOD LION	63fcd7cea4f8442c3386b589	7.833997e+11	1.00	3.49
Receip [.]	t ID: 00026b4c-dfe8-49dd-b026-4c2f0	fd5c6a1 PURCHASE_DATE	SCAN DATE	STORE_NAME	USER ID	PARCON	E FINAL QUANTITY	, EINIAI SAIE
4	00026b4c-dfe8-49dd-b026-4c2f0fd5c6a1		2024-07-05 15:56:43.549 Z	RANDALLS	6193231ae9b3d75037b0f928			
43233	00026b4c-dfe8-49dd-b026-4c2f0fd5c6a1		2024-07-05 15:56:43.549 Z	-	6193231ae9b3d75037b0f928			
43233	0002004C-0160-4900-0020-4C210103C0d1	2024-07-04	2024-07-03 13.30.43.349 2	RAINDALLS	01932314e9b3073037b01920	4.790050e+10	J 1.00	5.29
Receip	t ID: 0002d8cd-1701-4cdd-a524-b7040	2e2dbc0						
	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_I	D BARCOI	DE FINAL_QUANTIT	Y FINAL_SALE
5	0002d8cd-1701-4cdd-a524-b70402e2dbc0	2024-06-24	2024-06-24 19:44:54.247 Z	Z WALMART	T 5dcc6c510040a012b8e7692	4 6.811314e+	11 ze	ro 1.46
40388	0002d8cd-1701-4cdd-a524-b70402e2dbc0	2024-06-24	2024-06-24 19:44:54.247 Z	Z WALMART	T 5dcc6c510040a012b8e7692	4 6.811314e+	11 1.0	00 1.46
Receip [.]	t ID: 000550b2-1480-4c07-950f-ff601	f242152						
	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FINAL_QUANTITY	FINAL_SALE
6	000550b2-1480-4c07-950f-ff601f242152	2024-07-06	2024-07-06 19:27:48.586 Z	WALMART	5f850bc9cf9431165f3ac175	4.920091e+10	1.00	
47862	000550b2-1480-4c07-950f-ff601f242152	2024-07-06	2024-07-06 19:27:48.586 Z	WALMART	5f850bc9cf9431165f3ac175	4.920091e+10	1.00	3.12
Receip [.]	t ID: 00096c49-8b04-42f9-88ce-941c5			CTOPE WITE		B45655	PINIAL ALIAN	PINIAL CA:-
		PURCHASE_DATE		STORE_NAME	USER_ID		FINAL_QUANTITY	
	00096c49-8b04-42f9-88ce-941c5e06c4a7		2024-08-21 17:35:21.902 Z		6144f4f1f3ef696919f54b5c		zero	3.59
36036	00096c49-8b04-42f9-88ce-941c5e06c4a7	2024-08-19	2024-08-21 17:35:21.902 Z	TARGET	6144f4f1f3ef696919f54b5c	7.830007e+10	1.00	3.59

Receipt ID: 000e1d35-15e5-46c6-b6b3-33653ed3d27e

	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FINAL_QUANTITY	FINAL_SALE	
8	000e1d35-15e5-46c6-b6b3-33653ed3d27e	2024-08-13	2024-08-13 18:21:07.931 Z	WALMART	61a6d926f998e47aad33db66	5.200001e+10	1.00		
41970	000e1d35-15e5-46c6-b6b3-33653ed3d27e	2024-08-13	2024-08-13 18:21:07.931 Z	WALMART	61a6d926f998e47aad33db66	5.200001e+10	1.00	0.98	

Receipt ID: 0010d87d-1ad2-4e5e-9a25-cec736919d15

	RECEIPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FINAL_QUANTITY	FINAL_SALE
9	0010d87d-1ad2-4e5e-9a25-cec736919d15	2024-08-04	2024-08-04 18:01:47.787 Z	ALDI	66686fc2e04f743a096ea808	NaN	zero	2.29
40976	0010d87d-1ad2-4e5e-9a25-cec736919d15	2024-08-04	2024-08-04 18:01:47.787 Z	ALDI	66686fc2e04f743a096ea808	NaN	1.00	2.29

ii Insights

1. Duplicate Entries in the Transactions Table

The Transactions Table contains duplicate entries for every unique receipt or transaction in the dataset, with occurrences of **2**, **4**, **6**, **8**, **or 12 duplicates** per transaction.

2. Inconsistencies in Duplicate Transactions with Same RECEIPT_ID

For transactions with the same **RECEIPT_ID**, there are discrepancies in fields like FINAL_QUANTITY and FINAL_SALE:

- In some cases, the FINAL_QUANTITY field differs between duplicates, with one entry showing the value "zero" (as a string), while another records are numeric value.
- Similarly, the FINAL_SALE field varies, with some entries showing a valid sale amount, while others display NA.

2. Missing Values

Users Dataset

ii Insights

- BIRTH_DATE: 3,675 missing values.
- STATE: 4,812 missing values, which accounts for a significant portion of users.
- LANGUAGE: 30,508 missing values—a large gap indicating issues with language preference collection.
- GENDER: 5,892 missing values.

Products Dataset

```
Out[7]: CATEGORY_1 111
CATEGORY_2 1424
CATEGORY_3 60566
CATEGORY_4 778093
MANUFACTURER 226474
BRAND 226472
BARCODE 4025
dtype: int64
```

ii Insights

- BARCODE: 4,025 missing values (products without unique identifiers).
- CATEGORY_1 to CATEGORY_4: Increasing missingness as we move deeper into categories (e.g., 778,093 missing in CATEGORY_4).
- MANUFACTURER and BRAND : Both missing for 226,474 products

The lack of barcodes and category data makes it impossible to uniquely identify a perticular product.

Transactions Dataset

ii Insights

• BARCODE: 5,762 missing values

The lack of barcodes in transactions data prevents linking a significant portion of transactions to products.

'DE' 'MI' 'IL' 'MS' 'WA' 'KS' 'CT' 'OR' 'UT' 'MD' 'OK' 'NE' 'NV' 'AL'

'AK' 'AR' 'HI' 'ME' 'ND' 'ID' 'WY' 'MT' 'SD' 'VT']

3. Inconsistent Data

STATE Field (Users Dataset)

```
In [9]: # Get a List of unique states in the users dataset
unique_states = users[users['STATE'].notna()]['STATE'].unique()

# Number of unique states
print(f"No of Unique States in the Dataset: {len(unique_states)}")

# List of unique states
print(users[users['STATE'].notna()]['STATE'].unique())

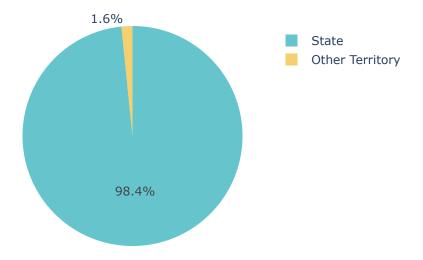
No of Unique States in the Dataset: STATE'].unique())

No of Unique States in the Dataset: STATE'].unique())

No of Unique States in the Dataset: STATE'].unique())

No of Unique States in the Dataset: STATE'].unique())
```

Distribution of States and Other Territories



ii Insights

• The STATE attribute contains 52 unique values, including DC (District of Columbia) and PR (Puerto Rico), which are not U.S. states.

2024-07-29 2024-07-29 21:06:17.211 Z

WALGREENS

DATE Fields (Transaction Dataset)

2213 16c9c13c-df14-4504-b802-108742e6a618

```
In [11]: txns[['RECEIPT_ID', 'STORE_NAME', 'PURCHASE_DATE', 'SCAN_DATE']].sample(2)

Out[11]: RECEIPT_ID STORE_NAME PURCHASE_DATE SCAN_DATE

5975 3d057fb4-9e30-48e9-b90f-0d8dee2b2e11 COSTCO 2024-06-24 2024-07-04 10:09:09.201 Z
```

Insights

• The PURCHASE_DATE field uses the MM/DD/YYYY format, while the SCAN_DATE field follows the ISO-8601 format with a timestamp (e.g., YYYY-MM-DDTHH:MM:SSZ).

GENDER Field (Users Table)

```
In [12]: users['GENDER'].value_counts()
Out[12]: GENDER
                                  64240
         female
         male
                                  25829
                                  1772
         transgender
         prefer_not_to_say
                                  1350
         non_binary
                                   473
                                   196
         unknown
                                   180
         not_listed
         Non-Binary
                                    34
                                    28
         not_specified
         My gender isn't listed
                                     5
         Prefer not to say
         Name: count, dtype: int64
         Insights
```

• Inconsistent variations exist for the same categories, such as:

- Non-Binary → Non-Binary , non_binary , not_listed , and "My gender isn't listed" .
- Prefer not to say → Prefer not to say , prefer_not_to_say .


4. Outliers

SCAN_DATE & PURCHASE_DATE Fields (Transactions Dataset)

```
In [13]: txns_cpy = txns.copy()
         # Convert dates to datetime format
          txns_cpy['PURCHASE_DATE'] = pd.to_datetime(txns_cpy['PURCHASE_DATE'], format='mixed')
          txns_cpy['SCAN_DATE'] = pd.to_datetime(txns_cpy['SCAN_DATE'])
          # Remove timezone from SCAN_DATE to make it same as PURCHASE_DATE
          txns_cpy['SCAN_DATE'] = pd.to_datetime(txns['SCAN_DATE']).dt.tz_localize(None)
         # Calculate the difference in days
          txns_cpy['DATE_DIFFERENCE'] = (txns_cpy['SCAN_DATE'] - txns_cpy['PURCHASE_DATE']).dt.days
         # Categorize the transactions
         def categorize_scan(diff):
             if diff < 0:</pre>
                 return 'Before Purchase'
             elif diff == 0:
                 return 'Same Day as Purchase'
             else:
                 return 'After Purchase'
          txns_cpy['SCAN_CATEGORY'] = txns_cpy['DATE_DIFFERENCE'].apply(categorize_scan)
```

```
# Count the occurrences of each category
scan_counts = txns_cpy['SCAN_CATEGORY'].value_counts().reset_index()
scan_counts.columns = ['Scan Category', 'Count']
# Create a pie chart
fig = px.pie(scan_counts,
            names='Scan Category',
            values='Count',
            title='Scan Timing Distribution Compared to Purchase Date',
            color_discrete_sequence=px.colors.qualitative.Pastel,
            width=500, # Adjust width to tighten the figure
            height=400 # Adjust height accordingly
# show count and percentage both in labels
fig.update_traces(textinfo='percent+value',
                 texttemplate='%{value} (%{percent})',
                 insidetextorientation='radial')
# Display the pie chart
fig.show()
```

Scan Timing Distribution Compared to Purchase Date



ii Insights


• There are **94 records** (0.188% of the dataset) where the SCAN_DATE occurs **before** the PURCHASE_DATE. This inconsistency is unusual, as a product is generally expected to be scanned only after it is purchased.

FINAL_QUANTITY Field (Transactions Dataset)

```
title="Scatter Plot of FINAL_QUANTITY with Outliers",
    labels={'FINAL_QUANTITY': 'Quantity'},
    color_discrete_sequence=px.colors.qualitative.Pastel,
    points='all',
    width=600, # Adjust width to tighten the figure
    height=400 # Adjust height accordingly
)

fig.show()
```

Scatter Plot of FINAL_QUANTITY with Outliers



```
In [15]: # Display Transactions with Quanity 276 (i.e Outliers)
txns[txns['FINAL_QUANTITY']=='276.00'][["RECEIPT_ID", "STORE_NAME", "USER_ID", "BARCODE", "FINAL_QUANTITY", "FINAL_SALE"]]

Out[15]: RECEIPT_ID STORE_NAME USER_ID BARCODE FINAL_QUANTITY FINAL_SALE
```

24810	fe0780d1-2d02-4822-8f12-7056b1814f17	MAIN STREET MARKET	5d197f9dd08976510c49d0e6	4.800135e+10	276.00	
42410	fe0780d1-2d02-4822-8f12-7056b1814f17	MAIN STREET MARKET	5d197f9dd08976510c49d0e6	4.800135e+10	276.00	5.89

ii Insights

• There are **two identical transactions** at **MAIN STREET MARKET** with unusually high FINAL_QUANTITY values of **276**, which appear to be potential **outliers**.

Q2: Are there any fields that are challenging to understand?

```
In [16]: # Display transactions with "zero" FINAL_QUANTITY
txns[txns['FINAL_QUANTITY']=="zero"].head(2)
```

Out[16]:	RECE	IPT_ID	PURCHASE_DATE	SCAN_DATE	STORE_NAME	USER_ID	BARCODE	FINAL_QUANTITY	FINAL_SALE
	1 0001455d-7a92-4a7b-a1d2-c747af	f1c8fd3	2024-07-20	2024-07-20 09:50:24.206 Z	ALDI	62c08877baa38d1a1f6c211a	NaN	zero	1.49
	3 000239aa-3478-453d-801e-66a82e	e39c8af	2024-06-18	2024-06-19 11:03:37.468 Z	FOOD LION	63fcd7cea4f8442c3386b589	7.833997e+11	zero	3.49
In [17]:	users[['ID','LANGUAGE']].head()								
Out[17]:	ID LAN	NGUAGE	_						
	0 5ef3b4f17053ab141787697d	es-419							
	1 5ff220d383fcfc12622b96bc	en							
	2 6477950aa55bb77a0e27ee10	es-419							
	3 658a306e99b40f103b63ccf8	en							

Insights

4 653cf5d6a225ea102b7ecdc2

Some fields in the dataset have logical inconsistencies that make them challenging to understand. For example:

- Transactions where FINAL_QUANTITY = 0 but FINAL_SALE is non-zero are illogical, as a non-zero sale amount should typically correspond to a positive quantity.
- The LANGUAGE field in USERS dataset contains codes such as en , es-419 , but the context is unclear. It is not clear if this represents the preferred language of users, the language in which receipts are submitted, or something else.

One Interesting Insights

Month-over-Month (MoM) growth of active user

```
In [18]: txns['SCAN_DATE'] = pd.to_datetime(txns['SCAN_DATE'], format='mixed')

# Extract Year and Month from PURCHASE_DATE
txns['YEAR_MONTH'] = txns['SCAN_DATE'].dt.to_period('M'))

# Group by YEAR_MONTH and count distinct active users
active_users_per_month = (
    txns.groupby('YEAR_MONTH')['USER_ID']
    .nunique()
    .reset_index()
    .reset_index()
    .rename(columns=('USER_ID': 'ACTIVE_USERS'))
)

# Calculate Month-over-Month (Mon) Growth
active_users_per_month['PREVIOUS_MONTH_USERS'] = active_users_per_month['ACTIVE_USERS'].shift(1)
active_users_per_month['Mon_GROWTH_PERCENTAGE'] = (
    (active_users_per_month['Mon_GROWTH_PERCENTAGE'] = (
    (active_users_per_month['ACTIVE_USERS'] - active_users_per_month['PREVIOUS_MONTH_USERS'])
    / active_users_per_month['PREVIOUS_MONTH_USERS'] * 100
)

# Drop rows without growth data (first month will have NaN for growth percentage)
```

```
# active_users_per_month = active_users_per_month.dropna(subset=['MOM_GROWTH_PERCENTAGE'])

# Sort results by most recent month
active_users_per_month = active_users_per_month.sort_values(by='YEAR_MONTH')

active_users_per_month
```

Out[18]: YEAR MONTH ACTIVE USERS PREVIOUS MONTH USERS MOM GROWTH PERCENTAGE

0	2024-06	4334	NaN	NaN
1	2024-07	8063	4334.0	86.040609
2	2024-08	7440	8063.0	-7.726653
3	2024-09	2092	7440.0	-71.881720

```
In [19]: # Convert YEAR_MONTH from Period to timestamp and then to string for visualization
         active_users_per_month['YEAR_MONTH'] = active_users_per_month['YEAR_MONTH'].dt.to_timestamp().dt.strftime('%Y-%m')
         # Create figure with dual y-axes
         fig = go.Figure()
         # Add bar chart for total active users on left y-axis
         fig.add_trace(go.Bar(
             x=active_users_per_month['YEAR_MONTH'],
             y=active_users_per_month['ACTIVE_USERS'],
             name='Active Users',
             marker_color=px.colors.qualitative.Pastel[1],
             text=active_users_per_month['ACTIVE_USERS'],
             textposition='outside',
             yaxis='y1' # Assign to primary y-axis (left)
         # Add line chart for MoM growth percentage on right y-axis
         fig.add_trace(go.Scatter(
             x=active_users_per_month['YEAR_MONTH'],
             y=active_users_per_month['MOM_GROWTH_PERCENTAGE'],
             name='MoM Growth (%)',
             mode='lines+markers+text',
             text=[f'{x:.1f}%' for x in active_users_per_month['MOM_GROWTH_PERCENTAGE']],
             textposition='bottom left',
             line=dict(color=px.colors.qualitative.Pastel[6], width=3),
             yaxis='y2' # Assign to secondary y-axis (right)
         ))
         # Add labels and format the graph
         fig.update_layout(
             title='Month-over-Month Growth of Active Users',
             height = 600,
             width = 600,
             xaxis=dict(
                 title='Month',
                 type='category',
                 categoryorder='array',
                 categoryarray=active_users_per_month['YEAR_MONTH'].tolist()
             ),
             yaxis=dict(
                 title='Active Users',
                 side='left',
                 showgrid=False
```

```
),
yaxis2=dict(
    title='MoM Growth (%)',
    side='right',
    overlaying='y',
    showgrid=False,
    ticksuffix='%',
    zeroline=False
),
legend=dict(title='Metrics'),
    template='plotly_white'
)
fig.show()
```

Month-over-Month Growth of Active Users



ii Insights

1. Strong Growth in July:

• Active users surged by **86%**, from **4,334** to **8,063**

2. Moderate Decline in August:

• User count dropped by **7.7%**

3. Sharp Drop in September:

• A significant **71.9%** decline to **2,092** users